From One Newsletter to a Curation Platform

Productizing the engine behind *The Intersect*

A discussion paper · Juergen Berkessel · Polymash · April 2026

THE QUESTION

Is The Intersect Curator a product, or a one-off?

WHAT EXISTS

A single-tenant Next.js app that ingests 200+ RSS feeds, scores with a Pinecone taste profile + Claude, assembles a newsletter, and produces a podcast end-to-end. In production weekly.

WHAT'S POSSIBLE

The same engine, configured for any two-topic intersection — real estate × rates, EV × supply chain, biotech × FDA policy. Five thin seams per client, one shared core.

The engine works. The question is whether it's worth running more than one beat on it.

THE THESIS, IN ONE SLIDE

90% Generic · 10% Per-Client

SHARED CORE · 90%

- RSS ingestion with dedup
- Pinecone taste-profile vectors
- Airtable data layer
- Two-layer AI scoring loop
- Issue assembly + entry writing
- WonderCraft → Auphonic → Captivate
- Cover art, publishing, analytics

PER-CLIENT SEAMS · 10%

- RSS feed list (config)
- The intersection rubric (prompt)
- Taste profile (Pinecone namespace)
- Brand voice (newsletter + podcast)
- Visual identity (palette, covers)

The code doesn't know what domain it's in. That's the unlock.

THE ENGINE TODAY

End-to-End Pipeline

INGEST
200+ RSS feeds every 3 hours

→

DEDUP
Hash + title similarity

→

EMBED
OpenAI → Pinecone

→

SCORE
Haiku + terms + taste sim

CURATE
Human rates on mobile

→

ASSEMBLE
Theme + order + entries

→

NEWSLETTER
Curated.co delivery

→

PODCAST
WonderCraft → Auphonic → Captivate

Every box above is domain-agnostic. The rubric is the only thing that knows it's art × tech.

What Changes Per Client

1. **RSS feed list**
   30–50 URLs in an Airtable table. Pure curation work by a domain expert.

2. **Intersection rubric**
   Scoring prompt: axis definitions, calibration examples, named failure modes. The hardest artifact to produce.

3. **Taste profile**
   Per-client Pinecone namespace. Seeded from editor ratings. Improves with use.

4. **Brand voice**
   Newsletter-writer prompt + podcast host personas (WonderCraft voice IDs).

5. **Visual identity**
   Cover-art prompt template, palette, logo. Lightest-weight seam.

ARCHITECTURE OPTIONS

Tenant Isolation — Two Paths

RECOMMENDED

Option A · Base per tenant

One Airtable base + one Pinecone namespace per client. Central "tenants" config table resolves request → tenant context.

+ Hard isolation, simple backups, per-client customization
+ Matches existing `lib/airtable/client.ts` pattern
– Onboarding copies a template base

ALTERNATIVE

Option B · Shared base + client_id

One base, every row carries a `client_id`. Every query filters by it.

+ Easier cross-tenant reporting
– One missed filter = data leak
– Airtable has no row-level security
– Per-tenant backup/restore is painful

Go with A. Airtable is a *collaboration tool* — clients should see their base, nothing else.

THREE EXAMPLE VERTICALS

The Rubric Transfers. The Examples Don't.

Residential RE × Rates

Feeds: Calculated Risk, Fed FEDS Notes, Urban Institute, Bloomberg RE

HIGH: "2–10 spread and housing volume"
LOW: Generic crash content, agent-bait listings

EV × Supply Chain

Feeds: Benchmark Mineral, IEA, USGS, DOE LPO, *Inside EVs*

HIGH: "Indonesia nickel ban → Ford F-150 Lightning sourcing"
LOW: Pure Tesla results, pure lithium price moves

Biotech × FDA Policy

Feeds: Endpoints, STAT, *In the Pipeline*, FDA AdCom calendar

HIGH: "Accelerated approval hinged on new biomarker guidance"
LOW: Pure trial roundups, pure policy debate

Every beat has its own "fake intersection" failure mode. Naming them is what makes the rubric sharp.

A SECOND AXIS FAMILY

Topic × Geography — One Rubric, Many Editions

NATIONAL + METRO FAN-OUT

Public Art × Urban Policy

National core: percent-for-art programs, public commissioning, public-art law.

Fan out: LA, NYC, Chicago editions — same rubric; feeds and scoring filter to each metro's commissioning agencies, mural permitting, and local panels.

One editorial investment → N monetizable editions, each with its own local sponsor market.

PLACE IS THE AXIS

Regional Craft Economy × Western NC

The two axes: a sector (craft, makers, studios) and a named region. Feeds scope to state craft councils, local press, regional maker publications.

Rubric: scores topical fit *and* geographic relevance. A national ceramics-week recap doesn't qualify without a regional tie-in.

Deep local authority. Defensible against national aggregators.

ENGINEERING EFFORT

4–8 Weeks to a 2-Tenant MVP

3–5 days	Tenant resolver middleware + request context
2–3 days	Airtable client → tenant-aware (base ID from context)
1–2 days	Pinecone → per-tenant namespaces
2–4 days	Service client key scoping (Curated, Captivate, WonderCraft)
5–10 days	Platform admin UI — onboarding, rubric editor, cold-start seeding
2–3 days	Per-tenant cron scheduling
3–5 days	Template base + automated tenant provisioning

Not included: rubric writing, feed curation, brand-voice capture. Those recur per client.

BUSINESS MODEL

Three Coherent Paths

Pure SaaS

$200–800/mo

Client configures & runs

Hard part: most buyers can't write a good rubric unaided

Done-for-You START HERE

$3k–8k/mo

Polymash delivers end-to-end

Hard part: scales linearly with headcount

Hybrid

$10k–25k setup

White-glove onboarding + platform fee

Hard part: needs both product *and* playbook

Start with done-for-you. Learn what the rubric workshop needs. Graduate when you know the playbook.

RISKS

What Could Kill This

Curated.co has no publish API

Manual click-to-send today. Fine at 3 clients. A part-time job at 5. Swap rail or live with it — decide early.

Voice licensing per client

Different clients want different hosts. Clone rights are a legal surface Polymash doesn't have a process for yet.

Rubric-writing is bottlenecked on Juergen

The thing that makes the product good doesn't scale past ~10 clients without a structured workshop format or a rubric-authoring assistant. Most important strategic question on this page.

NEXT STEPS

Cheap Validation Before Any Code

1 **Pick a real candidate vertical**

Not hypothetical. One where there's a plausible buyer intro in 30 days.

2 **Run a rubric workshop with the buyer**

Rubric + 20-article calibration set. Expert reads it, agrees with it.

3 **Swap the rubric into the existing engine locally**

Seed 50 articles. See how they rank. This kills 90% of the "does it work" risk before any tenant-isolation code.

4 **If step 3 passes → 2-tenant MVP in 6 weeks**

The Intersect itself + one new client. Scope as in the effort estimate.

5 **Pick a business model before pricing the engagement**

Onboarding fee structure depends entirely on which of the three paths we're on.

The Engine Works.

The question is whether Polymash wants to run it for more than one beat.

90%	5	6 wks
engine is generic	per-client seams	to 2-tenant MVP